CIS 520, Machine Learning, Fall 2018: Assignment 3
Due: Tuesday, October 2nd, 11:59pm

Instructions. Please write up your responses to the following problems clearly and concisely. We encourage you to write up your responses using LaTeX; we have provided a LaTeX template, available on Canvas, to make this easier. **Submit your answers in PDF form to Gradescope. We will not accept paper copies of the homework.**

Collaboration. You are allowed and encouraged to work together. You may discuss the homework to understand the problem and reach a solution in groups up to size **two students.** However, *each student must write down the solution independently, and without referring to written notes from the joint session.* **In addition, each student must write on the problem set the names of the people with whom you collaborated.** You must understand the solution well enough in order to reconstruct it by yourself. (This is for your own benefit: you have to take the exams alone.)

Collaborators:

Type Collaborator Name Here

1 Naïve Bayes as a Linear Classifier [25 points]

In this question we will consider the problem of binary classification, where we call one class positive and the other negative (for example spam vs. non-spam), i.e. each label $y \in \{\pm 1\}$. We will also assume that each instance $\mathbf{x} = (x_1, \cdots, x_n)$ has binary attribute/feature values, i.e. each attribute/feature $x_i \in \{0, 1\}$.

Let $p = \Pr(y = 1)$, $\alpha_i = \Pr(x_i = 1 | y = 1)$, and $\beta_i = \Pr(x_i = 1 | y = -1)$. We will assume that all the attributes of each instance $\mathbf{x}$ are conditionally independent given y. Formally,

$$\Pr(\mathbf{x}|y) = \Pi_{i=1}^n \Pr(x_i|y).$$

Recall that a Naïve Bayes classifier h [1] can be written as:

$$h(\mathbf{x}) = \underset{y \in \{\pm 1\}}{\operatorname{argmax}} \widehat{\Pr}(y|\mathbf{x}), \tag{1}$$

where the probability $\widehat{\Pr}(y|\mathbf{x})$ is estimated from data. For the above problem the Naïve Bayes classifier can be written in the form of a linear classifier, i.e. for some $\mathbf{w} \in \mathbb{R}^n$ and $b \in \mathbb{R}$

$$h(\mathbf{x}) = \operatorname{sign}(\mathbf{w}^\top \mathbf{x} + b),$$

---

[1] Assume $h(\mathbf{x}) = -1$ in case there is a tie.

where the sign function returns $+1$ when $\mathbf{w}^\top \mathbf{x} + b$ is positive, and -1 otherwise. In the problem you will be asked to find such a $\mathbf{w}$ and b.

1. **[2 points]** Show that the conditional probability of $\mathbf{x}$ given y can be written as:

$$\Pr(\mathbf{x}|y = 1) = \Pi_{i=1}^{n} \ \alpha_i^{x_i} \cdot (1 - \alpha_i)^{(1-x_i)},$$

   and

$$\Pr(\mathbf{x}|y = -1) = \Pi_{i=1}^{n} \ \beta_i^{x_i} \cdot (1 - \beta_i)^{(1-x_i)}.$$

2. **[8 points]** Given data $D = \{(\mathbf{x}_1, y_1), \cdots (\mathbf{x}_m, y_m)\}$ find the maximum likelihood estimates (MLE) of the parameters p, α_i, and β_i for each $i \in \{1, \cdots, n\}$. Call these estimates $\widehat{p}$, $\widehat{\alpha}_i$, and $\widehat{\beta}_i$, respectively.

3. **[3 points]** Let $\widehat{\Pr}(y|\mathbf{x})$ be the probability distribution of y given $\mathbf{x}$ corresponding to the MLE estimates $\widehat{p}$, $\widehat{\alpha}_i$'s, and $\widehat{\beta}_i$'s. Using Equation (1) show that $h(\mathbf{x})$ can be written as

$$h(\mathbf{x}) = \text{sign}\big(\widehat{\Pr}(1|\mathbf{x}) - \widehat{\Pr}(-1|\mathbf{x})\big). \tag{2}$$

4. **[12 points]** Using Bayes rule and the form of $\widehat{\Pr}(y|\mathbf{x})$ show that

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \mathbf{x} + b),$$

   and find the value of $\mathbf{w}$ and b.

   Hint: You need to take the log of both $\widehat{\Pr}(1|\mathbf{x})$ and $\widehat{\Pr}(-1|\mathbf{x})$ in Equation (2), and use the fact that log is an increasing function.

2 Multiclass Logistic Regression [30 points]

In this question, we will see how we can extend the logistic regression model from HW2 (which was used for binary classifiction) to multi-class classification. Let's say we have C different classes, and for a class j we have :

$$\mathbf{P}(Y = j \mid X = \mathbf{x}) = \frac{\exp\{\mathbf{w}_j^T x\}}{\sum_{k=1}^{C} \exp\{\mathbf{w}_k^T x\}} \quad \forall j \in \{1, 2, .., C\}$$

where as usual $\mathbf{x}$ is a vector of features, and $\mathbf{w}_j$ is the weight vector assigned to class j. Our objective is to estimate the weights using gradient ascent (just like we did last week), but this time there will not be any coding involved. We will also add a regularization term to the loss function to avoid overfitting.

1. **[15points]** Suppose that the training matrix is of dimensions $N \times P$, which is to say that you have N data points and each data point has P features. Write down the log likelihood, $L(\mathbf{w}_1, ..., \mathbf{w}_C)$. Now add a $L2$ regularization term. Please show all your steps and write a justification for each step.

2. **[7points]** Next, derive the expression for the j^{th} index in the vector gradient (i.e. partial derivative) $L(\mathbf{w}_1, ..., \mathbf{w}_C)$, with respect to $\mathbf{w}_j$.

3. **[4points]** Now, write down the update equation for weight vector $\mathbf{w}_j$, with η as the step size.

4. **[4points]** Will the sequence of consecutive weight vectors converge? If yes, to what? Why?

3   Feature Selection   [23 points]

We saw in class that one can use a variety of regularization penalties in linear regression.

$$\widehat{w} = \arg\min_w \quad \|Y - Xw\|_2^2 + \lambda\|w\|_p^p$$

Consider the three cases, $p = 0$, 1, and 2. (Where, to be precise, the exponent p isn't there for $p = 0$.) We want to know what effect these different penalties have on estimates of w.

Let's see this using a simple problem. Use the provided data (data.mat). Assume the constant term in the regression is zero, and assume $\lambda = 1$, except, of course, for question (1). You don't need to write code that solves these problems in their full generality; instead, feel free to use matlab to do the main calculations. The best way to search over parameter spaces is using the Matlab function $fminsearch$.(*Note:* If you are not familiar with this function, please see Matlab documentation.)

1. [**3 points**] If we assume that the response variable y is distributed according to $y \sim N(w \cdot x, \sigma^2)$, then what is the MLE estimate $\widehat{w}_{MLE}$ of w?

2. [**2 points**] Given $\lambda = 1$, what is $\widehat{w}$ for $p = 2$?

3. [**2 points**] Given $\lambda = 1$, what is $\widehat{w}$ for $p = 1$?

4. [**4 points**]   Given $\lambda = 1$, what is $\widehat{w}$ for $p = 0$? Note that since L0 norm is not a "real" norm, the penalty expression is a little different:

$$\widehat{w} = \arg\min_w \quad \|Y - Xw\|_2^2 + \lambda\|w\|_0$$

Also, for the L_0 norm, you will have to solve the (combinatorially many) cases where different components of w are set to zero, then add the L_0 penalty to each based on the number of features. There are 8 cases for 3 unknown w_i.

5. [**5 points**] Write a paragraph describing the relation between the estimates of w in the four cases (i.e. the four estimates of w from the first four parts of this question), explaining why that makes sense given the different penalties.

6. [**7 points**] When $\lambda > 0$, we make a trade-off between minimizing the sum of squared errors and the magnitude of $\widehat{w}$. In the following questions, we will explore this trade-off further. For the following, use the same data from data.mat.

   (a) [**1 point**] For the MLE estimate of w (as in 4.1), write down the value of the ratio

   $$\|\widehat{w}_{MLE}\|_2^2 \; / \; \|Y - X\widehat{w}_{MLE}\|_2^2.$$

   (b)   i. [**2 points**] Suppose the assumptions of linear regression are satisfied. Let's say that with N training samples (assume $N \gg P$, where P is the number of features), you compute $\widehat{w}_{MLE}$. Then let's say you do the same, this time with $2N$ training samples. How do you expect $\|Y - X\widehat{w}_{MLE}\|_2^2$ to change when going from N to $2N$ samples? When $N \gg P$, does this sum of squared errors for linear regression directly depend on the number of training samples?

ii. [**2 points**] Likewise, if you double the number of training samples, how do you expect $||\widehat{w}_{MLE}||_2^2$ to change? Does $||\widehat{w}_{MLE}||_2^2$ for linear regression directly depend on the number of training samples in the large-N limit?

(c) [**1 point**] Using any method (e.g. trial and error, random search, etc.), find a value of λ for which the estimate $\widehat{w}$ satisfies

$$0.8 < ||\widehat{w}||_2^2 \; / \; ||\widehat{w}_{MLE}||_2^2 < 0.9.$$

(d) [**1 point**] Using any method (e.g. trial and error, random search, etc.), find a value of λ for which the estimate $\widehat{w}$ satisfies

$$0.4 < ||\widehat{w}||_2^2 \; / \; ||\widehat{w}_{MLE}||_2^2 < 0.5.$$

4   MDL on a toy dataset   [22 points]

We provide a data set (train_data.mat, train_y.mat, test_data.mat, test_y.mat) generated from a particular model with $N = 64$. We want to estimate

$$\widehat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3$$

We want to use MDL to find the 'optimal' L_0-penalized model.

1. [**12 points**] Estimate the three linear regressions

   (We could actually try all possible subsets here, but instead we'll just try three.)

   $$y_1 = w_1 x_1$$
   $$y_2 = w_1 x_1 + w_2 x_2$$
   $$y_3 = w_1 x_1 + w_2 x_2 + w_3 x_3$$

   For each of the three cases, what is

   (a) the sum of square error
   i) $\text{Err}_1 =$
   ii) $\text{Err}_2 =$
   iii) $\text{Err}_3 =$

   (b) 2 times the estimated bits to code the residual $(n \log \frac{Error}{n})$
   i) $\text{ERR\_bits}_1 =$
   ii) $\text{ERR\_bits}_2 =$
   iii) $\text{ERR\_bits}_3 =$

(c) 2 times the estimated bits to code each residual plus model under AIC ($2 * 1$ bit to code each feature)

i) $\text{AIC\_bits}_1 =$

ii) $\text{AIC\_bits}_2 =$

iii) $\text{AIC\_bits}_3 =$

(d) 2 times the estimated bits to code each residual plus model under BIC ($2 * (1/2)log(n)$ bits to code each feature)

i) $\text{BIC\_bits}_1 =$

ii) $\text{BIC\_bits}_2 =$

iii) $\text{BIC\_bits}_3 =$

2. [**5 points**] Which model has the smallest minimum description length?

a) for AIC

b) for BIC

3. [**5 points**] Included in the kit is a test data set; does the error on the test set for the three models correspond to what is expected from MDLs? Please compute and show the test errors, and briefly explain in one sentence. Please use your model to obtain these values, and know that we will be checking them against values obtained using $fminsearch$.